

Real-Time ML Platform — Capgemini

Saketh Sai Nigam Kanduri — Software Engineer, Capgemini

■ Project Overview

Developed and maintained a production-grade Real-Time ML Platform that integrated Data Engineering, MLOps, and Backend microservices on AWS and Databricks. Enabled enterprise apps to consume ML predictions with low latency and high reliability.

■ System Architecture

Data Sources	Kafka, RDS, S3, Partner APIs
Ingestion & Processing	Databricks Autoloader, Delta Lake (Bronze → Silver → Gold)
Feature Store	Databricks Feature Store (Train/Serve parity)
Model Management	MLflow Tracking & Registry
Serving Layer	FastAPI microservices (Dockerized) on EKS & OpenShift
CI/CD	Jenkins & GitHub Actions (Tests, Security, Helm Deploy)
Monitoring	Grafana, Prometheus, Splunk (Drift, Latency, Errors)

■ Key Results

- Improved system uptime and scalability by **35%**
- Achieved **<150 ms** average prediction latency
- Reduced model downtime via proactive drift detection and automated rollback
- Deployed containerized ML APIs across Kubernetes and OpenShift clusters seamlessly

■ Technical Stack

Python, FastAPI, Databricks, MLflow, Delta Lake, Kafka, AWS (EKS, S3, DMS, IAM), Docker, Jenkins, GitHub Actions, Helm, Prometheus, Grafana, Splunk

■ End-to-End Workflow

- 1■■ Ingest data from Kafka & RDS → Delta Lake via Databricks Autoloader
- 2■■ Transform & validate data → Register in Feature Store
- 3■■ Collaborate with Data Scientists → Train and version models in MLflow
- 4■■ Deploy model APIs with FastAPI → Containerize & push to EKS/OpenShift
- 5■■ CI/CD automation (build, test, deploy) via Jenkins & GitHub Actions
- 6■■ Continuous monitoring with Grafana & Splunk → Trigger rollback on anomalies

■ Summary

Delivered a scalable, observable, and automated ML deployment ecosystem bridging Data Engineering and MLOps. Enabled faster releases, resilient model rollouts, and measurable business impact for enterprise clients.